

Regularization and Optimization

Aziz Temirkhanov

LAMBDA lab
Faculty of Computer Science
Higher School of Economics

March 4, 2024

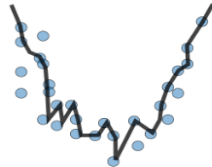

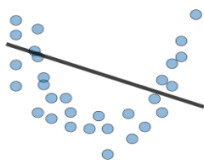
Outline

- 1 Overfitting
- 2 Dropout
- 3 Batch Normalization
- 4 Optimization

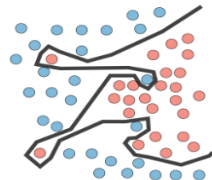

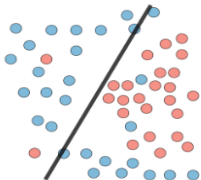
Overfitting

Overfitting


Regression
illustration



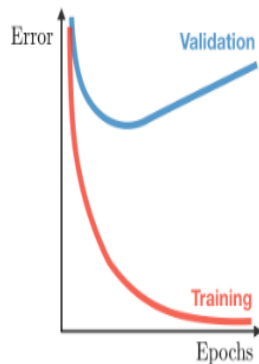
Classification
illustration



Overfitting



Overfitting



- Complexify model
- Add more features
- Train longer

- Perform regularization
- Get more data

Overfitting

Possible remedies:

- Add regularizing term to the loss function, e.g. penalize for complex parameters
- Introduce random noise to the input data
- Introduce random noise to the labels
- Acquire more data samples
- Change the network topology
- ...

Dropot

Dropout

Let $z \in \mathbb{R}^{B \times N}$, B – batch size, N – hidden size be a batch of latent representations, i.e. logit of some intermediate layer.

Sample a binary mask $m \in \{0, 1\}^{B \times N}$ $m_{ij} \sim \text{Bernoulli}(1 - p)$, where p – dropout rate, typically $0.1 \leq p \leq 0.5$

Dropout

Train mode:

- ① Sample new mask m
 - ② $y = m \odot z \times \frac{1}{1-p}$.
- Normalization by $\frac{1}{1-p}$ is required to preserve the expectation of neurons:

$$\mathbb{E}[y] = \mathbb{E}\left[m \odot z \cdot \frac{1}{1-p}\right] = \frac{1}{1-p} \cdot z \odot \mathbb{E}[m] = \frac{1}{1-p} \cdot z \cdot (1-p) = z \quad (1)$$

- The mask needs to be stored for backward pass

Dropout

Eval mode:

- As we normalized the output during training, it is an identical transform during testing:

$$y = z$$

- Dropout makes training slower but improves test quality

Batch Normalization

Batch Normalization

Main idea:

stabilize training by forcing zero mean and unit variance of features

Let:

$$X \in \mathbb{R}^{B \times N} = (x_1^{(N)}, x_2^{(N)}, \dots, x_B^{(N)})^T$$

Train mode:

1. Compute batch statistics:

$$\mu = \frac{1}{B} \sum_{i=1}^B x_i, \quad \sigma^2 = \frac{1}{B} \sum_{i=1}^B (x_i - \mu)^2 \quad (2)$$

$$\mu, \sigma^2 \in \mathbb{R}^N$$

2. Normalize x , recall that every operation is element-wise:

$$\hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 \epsilon}} \quad (3)$$

Batch Normalization

All operations are differentiable, so we need to compute

$$\frac{d\ell}{d\hat{x}_i}, \frac{d\ell}{d\mu}, \frac{d\ell}{d\sigma^2}, \frac{d\ell}{dx_i}$$

during backward pass

Then, apply trainable element-wise affine transform:

$$y_i = \hat{x}_i \odot w + b; \quad w, b \in \mathbb{R}^N \quad (4)$$

We give the model freedom to set the desired mean and variance for x_i

Batch Normalization

During the test, we do not have access to a batch of data.
Thus, we will accumulate statistics from the training batches:

$$\begin{aligned} running_mean &= running_mean \cdot (1 - m) + \mu \cdot m \\ running_var &= running_var \cdot (1 - m) + \sigma^2 \cdot m \cdot \frac{B}{B - 1} \end{aligned} \quad (5)$$

Here, m – is a momentum parameter (usually $m = 0.1$ to ensure slow updates).

In the beginning, *running_mean* is initialized with a zero-valued vector and *running_var* as ones. These vectors are not considered trainable.

Batch Normalization

- 1 Normalizing using accumulated statistics:

$$\hat{x}_i = \frac{x_i - \text{running_mean}}{\sqrt{\text{running_var} + \epsilon}} \quad (6)$$

- 2 Apply affine transform:

$$y_i = \hat{x}_i \odot w + b \quad (7)$$

Optimization

SGD

- $\{x_i, y_i\}_{i=1}^{\ell}$ – training set
- B – batch size
- $\mathcal{L}(\cdot, \cdot)$ – loss function
- W – all trainable parameters of the model
- $f_W(\cdot)$ – neural network

At the step t we have:

$$\mathcal{L}^t(w) = \frac{1}{B} \sum_{i=1}^B \mathbf{L}(f_W(x_i), y_i), \quad t_i \sim \mathcal{U}(\{1, \dots, \ell\}) \quad (8)$$

SGD

SGD algorithm:

① $g_t = \nabla_w \mathcal{L}^t(w_{t-1})$

② $m_t = \mu m_{t-1} + g_t$

③ $w_t = w_{t-1} - \eta_t m_t$

$m_0 = 0$, $0 \leq \mu \leq 1$, usually $\mu = 0.9$

SGD with Regularization

We can add L^2 **regularization**:

$$\mathcal{L}^t(w) \mapsto \mathcal{L}_\lambda^t(w) := \mathcal{L}^t(w) + \frac{\lambda}{2} \|w\|_2^2 \quad (9)$$

$$\nabla_w \|w\|_2^2 = 2w$$

$$\textcircled{1} \quad g_t = \nabla_w \mathcal{L}^t(w_{t-1}) + \lambda w_{t-1}$$

$$\textcircled{2} \quad m_t = \mu m_{t-1} + g_t$$

$$\textcircled{3} \quad w_t = w_{t-1} - \eta_t m_t$$

It is essentially a weight decaying by a factor of less than 1:

$$w_t = w_{t-1}(1 - \eta_t \lambda) - \eta_t(\mu m_{t-1} + \nabla_w \mathcal{L}^t(w_{t-1}))$$

Adam

- Adaptive learning rate (AdaFrad, RMSProp)
- Momentum

$$① \quad g_t = \nabla_w \mathcal{L}^t(w_{t-1}) + \lambda w_{t-1}$$

$$② \quad m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$③ \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 - \text{element-wise square}$$

$$④ \quad \hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$⑤ \quad w_t = w_{t-1} - \eta_t \frac{\hat{m}_t}{\hat{v}_t + \epsilon}$$

$m_0 = 0, v_0 = 0, \beta_1, \beta_2, \epsilon$ – hyperparams

$$\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$$

AdamW

- ① $g_t = \nabla_w \mathcal{L}^t(w_{t-1})$
- ② $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
- ③ $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$ – element-wise square
- ④ $\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$
- ⑤ $w_t = w_{t-1} (1 - \eta_t \lambda) - \eta_t \frac{\hat{m}_t}{\hat{v}_t + \epsilon}$

LR Schedulers

It is often a non-trivial task to choose a learning rate η_t

You can update it once per training epoch

- Constant. $\eta_t = \eta_0$. Often gives suboptimal quality
- StepLR. $\forall t \in \{t_0, \dots, t_n\} : \eta_{cur} = \eta_t, \eta_0 > \eta_1 > \dots$
- ExponentialLR. $\eta_t = \gamma \eta_{t-1}, \gamma < 1$
- CosineLR. $\eta_t = \eta_0 \cdot \frac{1}{2} (1 + \cos(\frac{\pi T}{T}))$, T – maximum number of epochs
- Linear warmup LR. Linearly increase lr first n -batches $\eta_0^{(0)} \mapsto \eta_0^{(n)}$.
Then apply any scheduler.
- Reduce LR on a plateau. Reduce LR when validation loss stagnates, thus requiring have validation set



Thank
you