

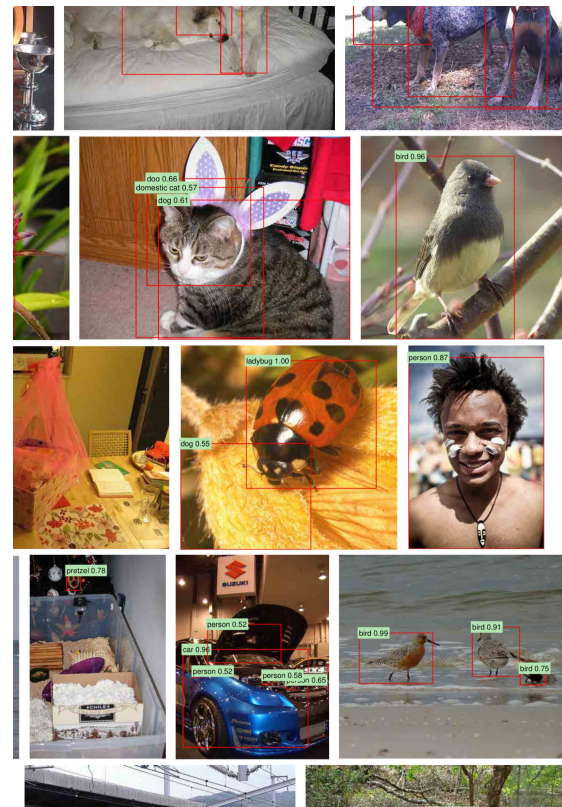
# Object Detection

---

Deep Learning

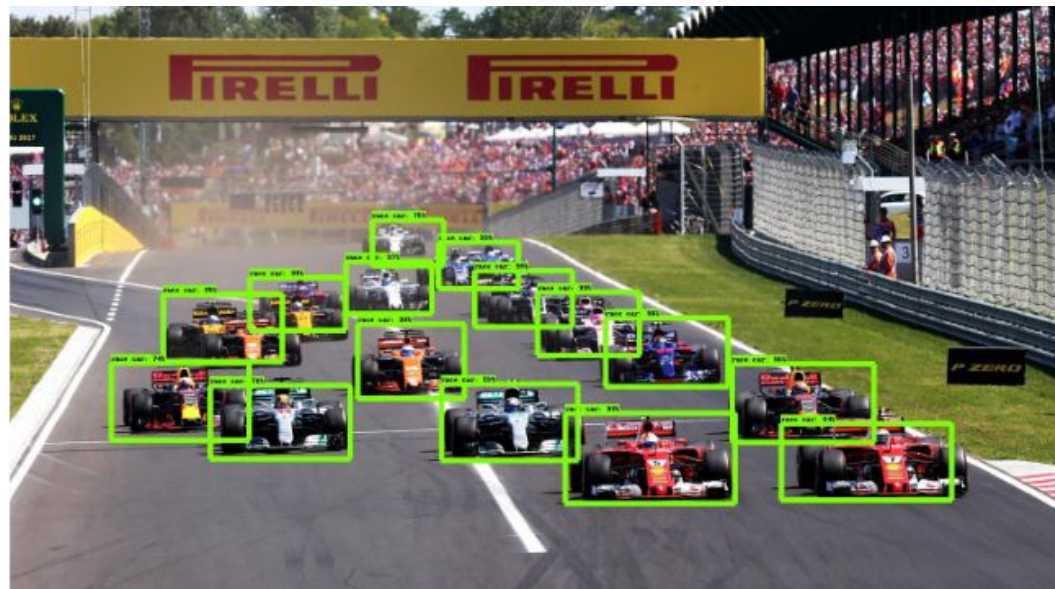
Aziz Temirkhanov  
Lambda, HSE

# Object Detection




# Object Detection

- **Goal:** find objects in the image and localize them with bounding boxes (BB)
- Variable number of objects in each image
- Intra-class variation of objects
- Imbalanced / rare classes
- Efficiency (FPS)




# Sliding Window

- Classify object and find coordinate of it
- Let's crop some region within the image and classify it
- The region with the highest probability of containing the object is the region we are looking for!
- In order to find best possible BB, iterate not only through the image, but also through window' hyperparams (e.g. window size)



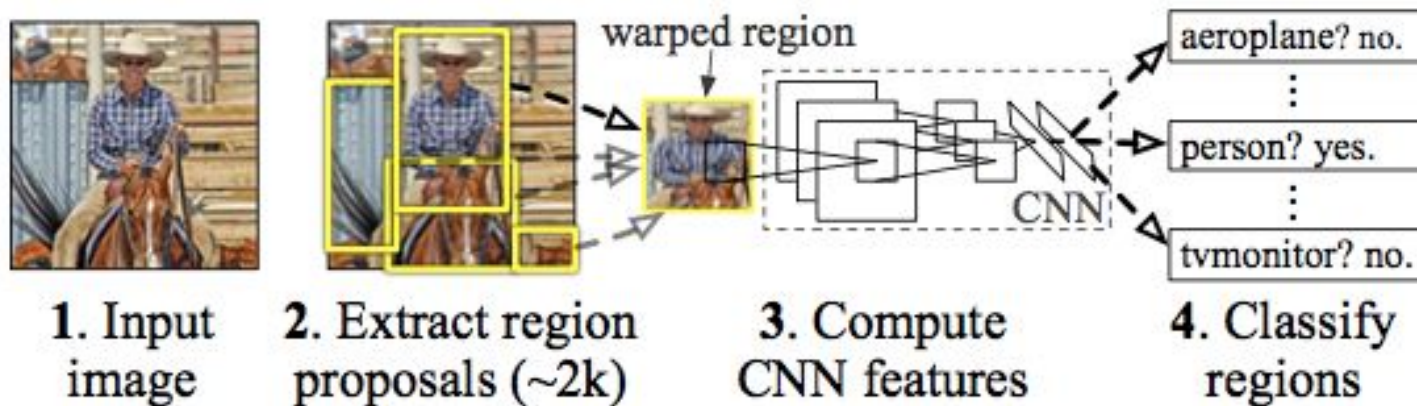
# Datasets

- Pascal VOC (Visual Object Classes), 20 classes
- MS COCO (Microsoft Common Objects in Context), 91 classes
- ImageNet, 200 classes



# Two-Stage Detectors

- Find regions with high probability of containing object
- Adjust BB in the second stage






# R-CNN


- Propose regions by Selective Search
  - Based on pixel intensity and using a graph-based algorithm, over-segment pixels
  - Group similar regions together
- Classify proposals and extract features using CNN
- Adjust coordinates
- Repeat for all proposals

Girshick et. al




# R-CNN

- Extract a vector from last layer of CNN (AlexNet in original work)
- $N_c + 1$  of total SVM classifiers solving one-vs-rest task
- Repeat for every proposal
- Adjust regions by  $N_c$  regressors
- Use feature maps

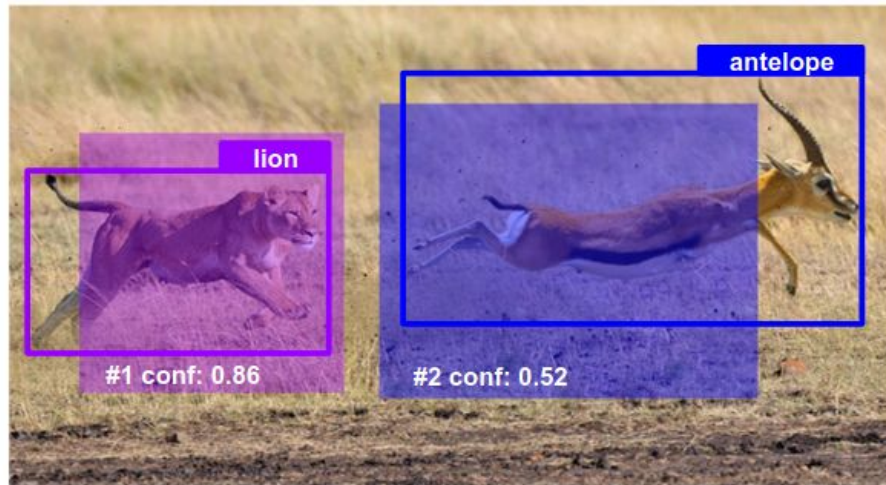




# IoU

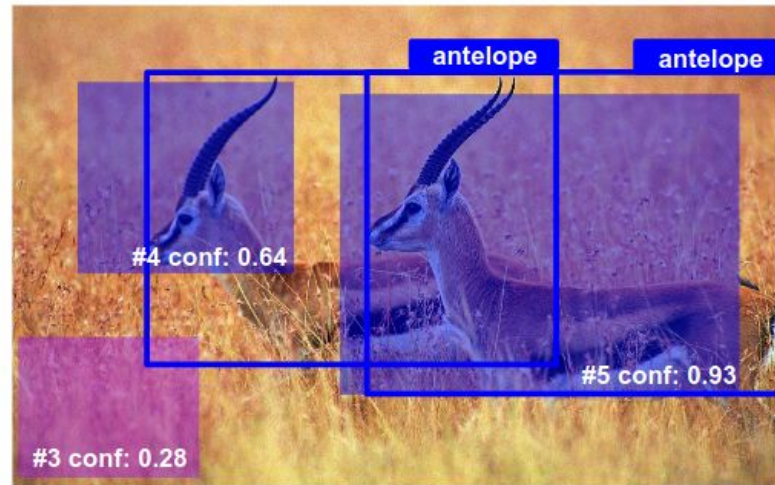


# mAP



Class: antelope

BB	conf	IoU (thr = 0.5)	TP/FP	precision	recall
#5	0.93	0.74	TP	1	0.33
#4	0.64	0.15	FP	0.5	0.33
#2	0.52	0.54	TP	0.66	0.66





Class: lion


BB	conf	IoU (thr = 0.5)	TP/FP	precision	recall
#1	0.86	0.68	TP	1	1
#3	0.28	0.0	FP	0.5	1

# mAP



- **AP** (average precision) = area under PR-curve
- **mAP** (mean AP) = AP averaged over all classes
- Different variants depending on IoU threshold:  $\text{mAP}_{50}$ ,  $\text{mAP}_{75}$ ,  $\text{mAP}_{[50:95]}$



...




...




IoU = 0.5

Loose



IoU = 0.7



IoU = 0.9


Tight

# Non Maximum Suppression

1. Take a set of BBoxes and sort them by certainty score
2. Select BBox with highest score and add it to final list of BBoxes, remove it from original list
3. Take next BBox in original list and compare it with BBoxes in selected list
4. If their IoU is higher than threshold remove second BBox

```
function nms(hypotheses, threshold):  
    sorted = sort(hypotheses.values, key=hypotheses.scores)  
    result = []  
    for first in sorted:  
        result.append(first)  
        without_first = sorted / first  
        for second in without_first:  
            if IoU(first, second) > threshold:  
                sorted.remove(second)  
    return result
```

Before non-maximum suppression



After non-maximum suppression



# Hard Negative Mining

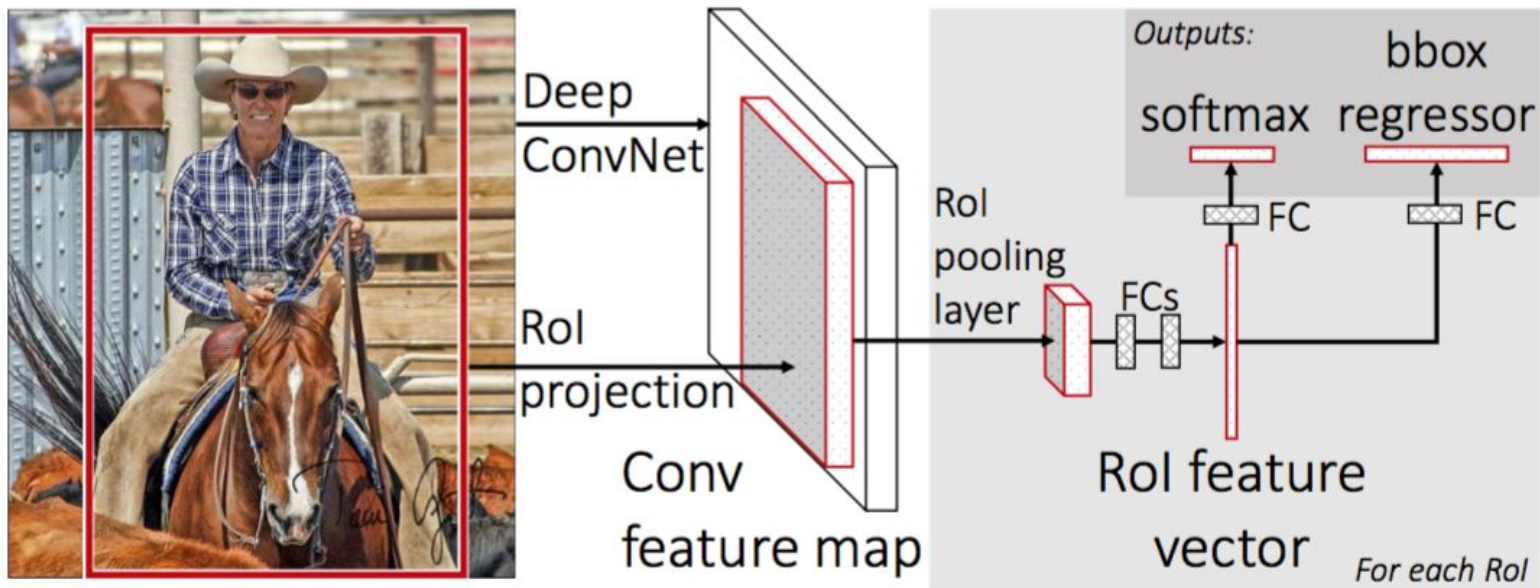
- Positive hypothesis — ones that contains an object
- Negative hypothesis — ones that contains a background or a part of an object
- Penalty classifier for false positive hp
- How to deal with Negative hypothesis?
  - Easy Negative — background
  - Hard Negative — wrong class or partial right class

# Hard Negative Mining

- Compare true BB with proposed one, using IoU
- Low IoU — easy negative
- High IoU — hard negative or positive
- Explicitly find those hard negative examples and add them to training data



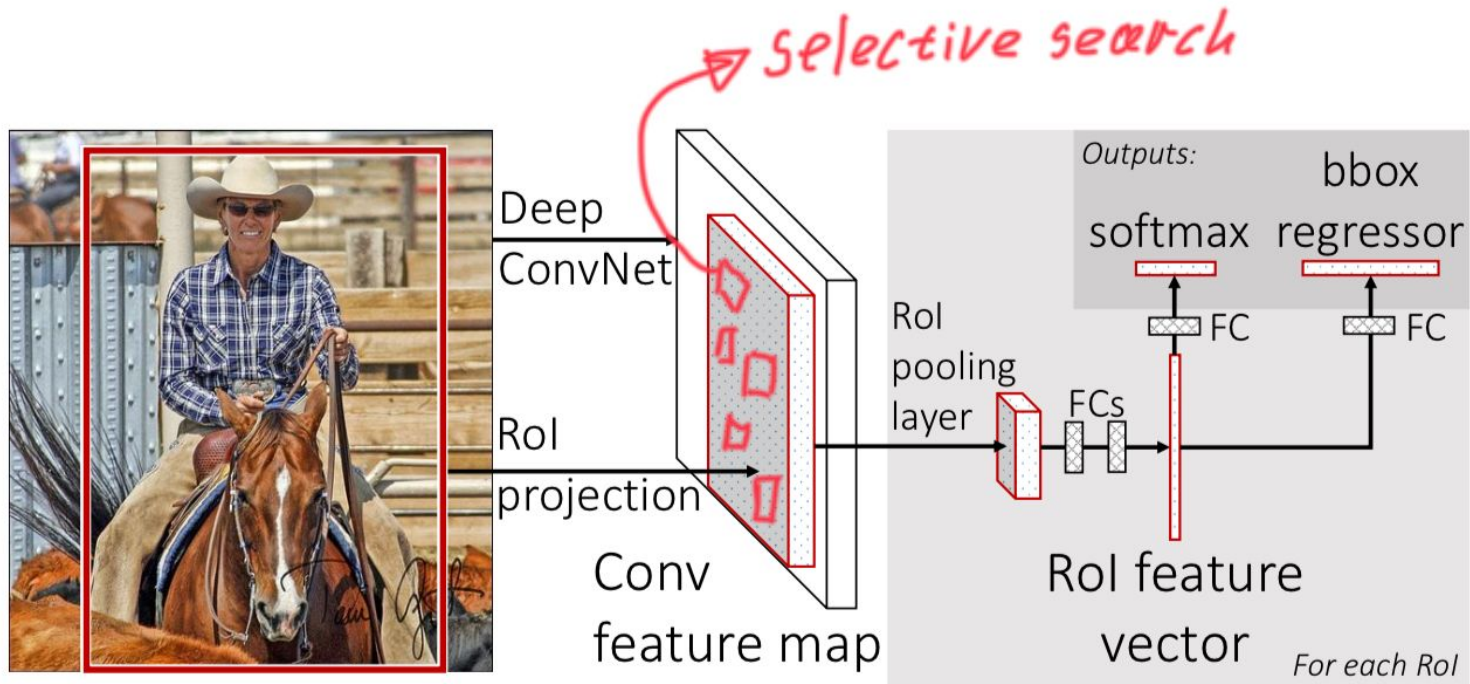
# Fast R-CNN



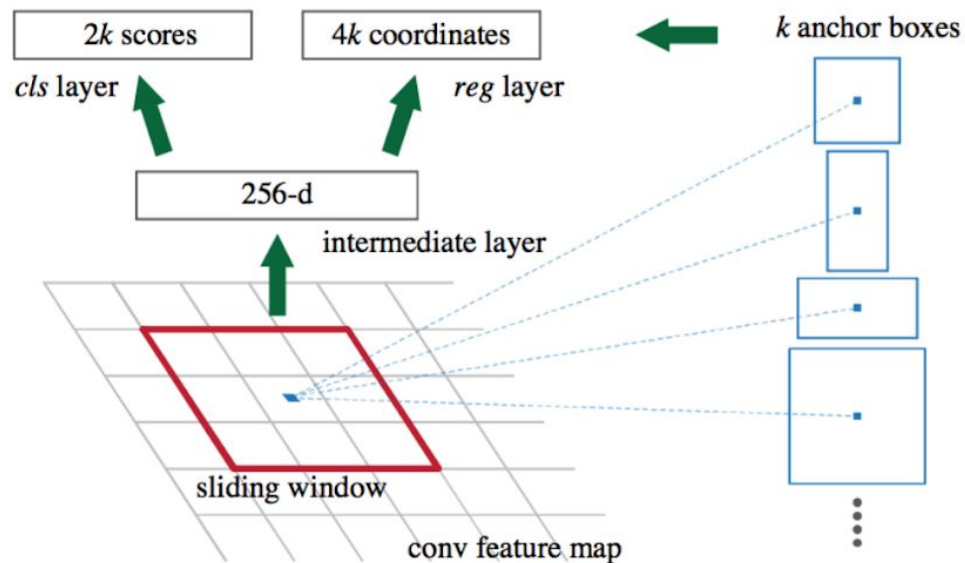
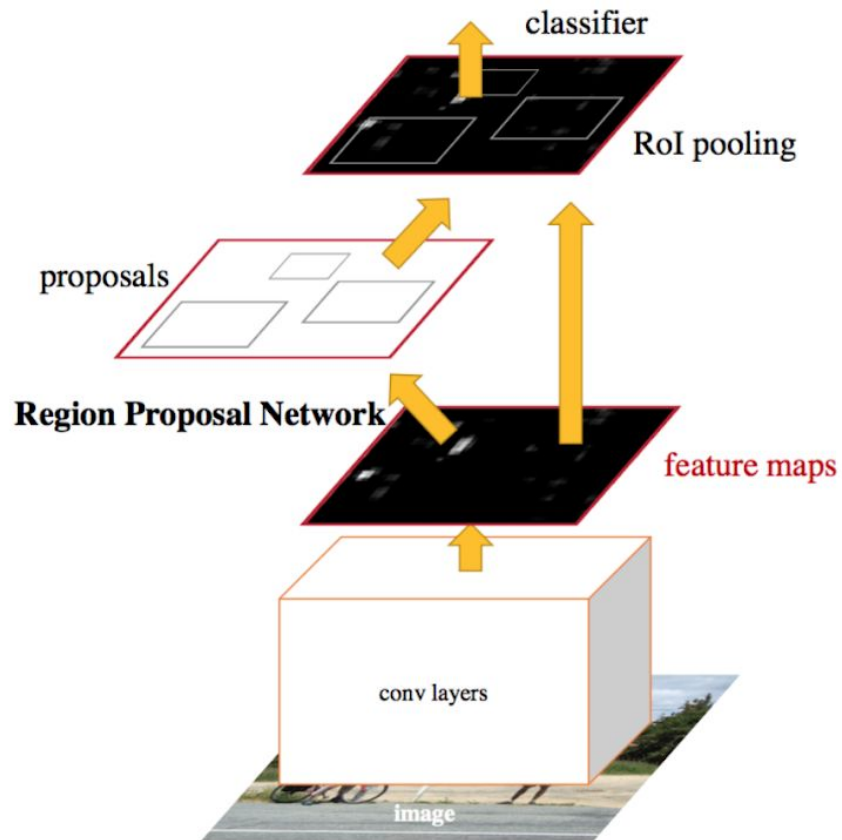
# Fast R-CNN

1. Feed image to a CNN block and obtain Feature Map
2. Run a Selective Search on the image and obtain proposals
3. Project proposal region onto feature map
4. Employ RoI pooling layer: adjust size of the region to be the same, and then maxpool the same-size filters

# Fast R-CNN



# Faster R-CNN



# Faster R-CNN

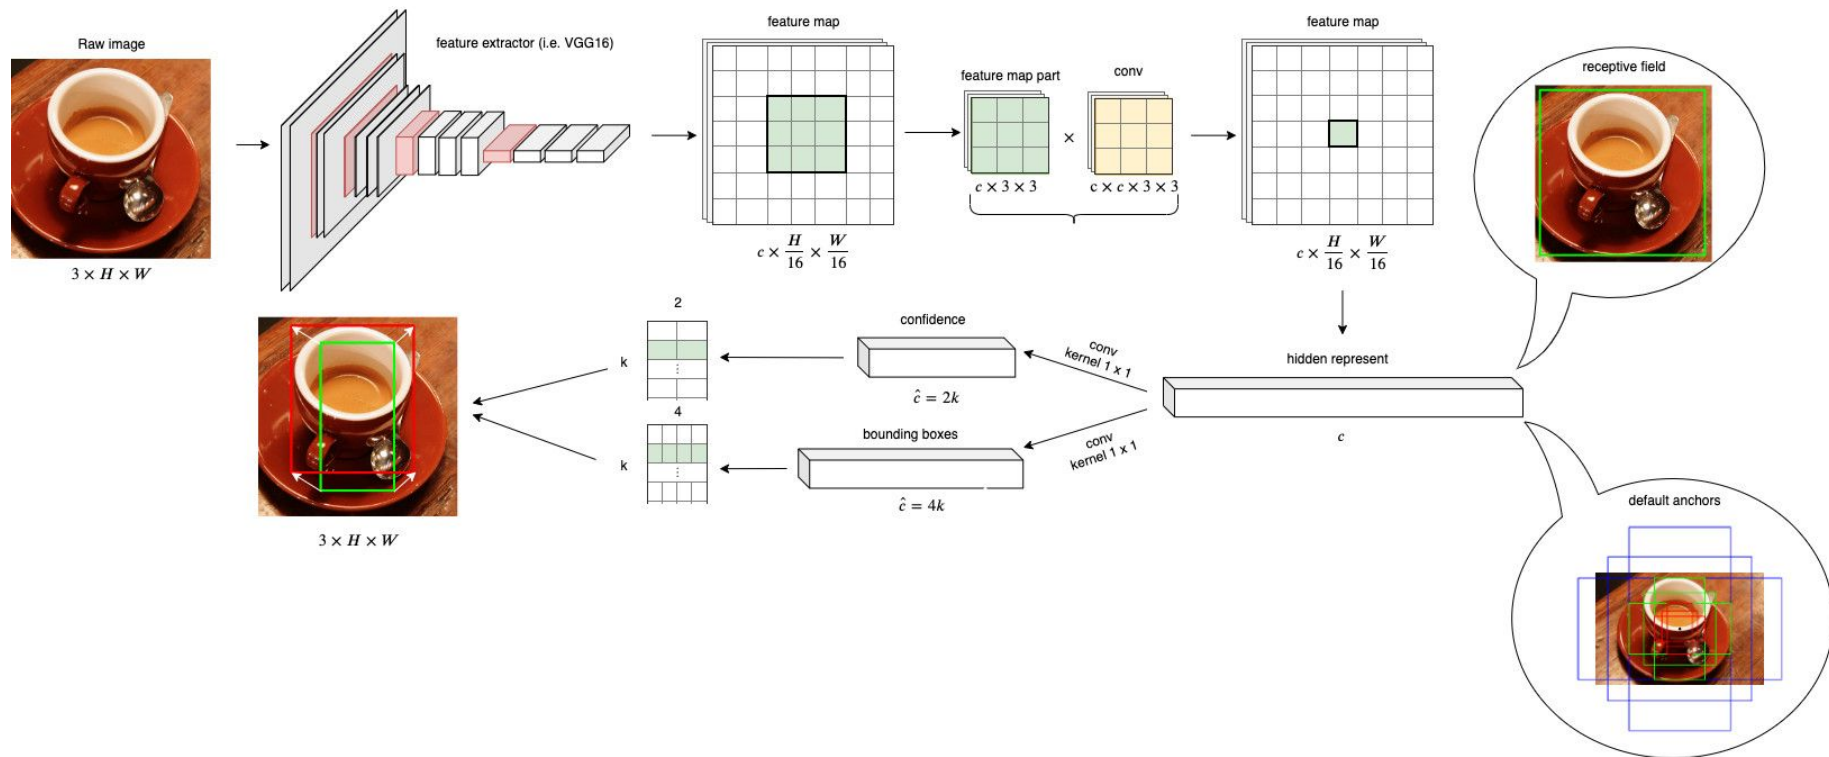
1. Pre-train a CNN network on image classification tasks.
2. Fine-tune the RPN (region proposal network) end-to-end for the region proposal task, which is initialized by the pre-train image classifier. Positive samples have IoU (intersection-over-union)  $> 0.7$ , while negative samples have IoU  $< 0.3$ .
  - a. Slide a small  $n \times n$  spatial window over the conv feature map of the entire image.
  - b. At the center of each sliding window, we predict multiple regions of various scales and ratios simultaneously. An anchor is a combination of (sliding window center, scale, ratio). For example, 3 scales + 3 ratios  $\Rightarrow k=9$  anchors at each sliding position.
3. Train a Fast R-CNN object detection model using the proposals generated by the current RPN
4. Then use the Fast R-CNN network to initialize RPN training. While keeping the shared convolutional layers, only fine-tune the RPN-specific layers. At this stage, RPN and the detection network have shared convolutional layers!
5. Finally fine-tune the unique layers of Fast R-CNN
6. Step 4-5 can be repeated to train RPN and Fast R-CNN alternatively if needed.

# Region Proposal Network

- Using VGG16' last conv layer — conv5\_3:
  - Effective stride: 16
  - Receptive Field size: 196
  - number of channels (feature maps): 512
- Take this feature maps and propose k hypothesis (k=9 at the beggining) for different sizes and aspect ratio. For default size it  $14 \times 14 \times 9 = 1764$  hypothesis
- Take  $c \frac{H}{16} \frac{W}{16}$  feature map and apply 3x3 convolution layer with stride=1
- Apply two 1x1 convolutions simultaneously to the feature map
  - cls layer translates 512 filters to 2k filters (binary classification for objects)
  - reg layer translates 512 filters to 4k filters: coordinates for each hypothesis
- Employ anchors: a defaults BBoxes with 3 different sizes (128x128, 256x256, 512x512) and 3 different aspect ratios (1:1, 2:1, 1:2), so 9 anchors in total.
- Now adjust this anchors with reg layer



# Faster R-CNN



# Faster R-CNN

1. Initialize RPN with CNN layer and generate current proposals
2. Train Fast R-CNN model using current proposals. Initialize weights with CNN
3. Train RPN network again using Fast R-CNN weights. At this stage, RPN and Fast R-CNN network shares the same layers, so freeze those layers and train only RPN-specific ones
4. Fine-tune Fast R-CNN network
5. Steps 3 and 4 can be repeated if needed

# One-Stage Detectors

---


Bonus part

## Two-stage and one-stage detectors


- Family of R-CNN models are all two-stage detectors, which can be slow
- One-Stage detectors runs once over dense sampling of possible locations

# YOLO

1. Pre-train the CNN model
2. Split image into  $S \times S$  cells
3. Assign  $B$  bounding boxes to each cell. Initial size is set via K-means and IoU throughout true bboxes
4. Predict 4 coordinates: x-center, y-center, h, w
5. Predict probability of containing object
6. Predict  $K$  conditional probabilities




$S \times S \times B$  bounding boxes  
 $\text{confidence} = Pr(\text{object}) \times \text{IoU}(\text{pred}, \text{truth})$



$S \times S$  grid on input




Bounding boxes + confidence



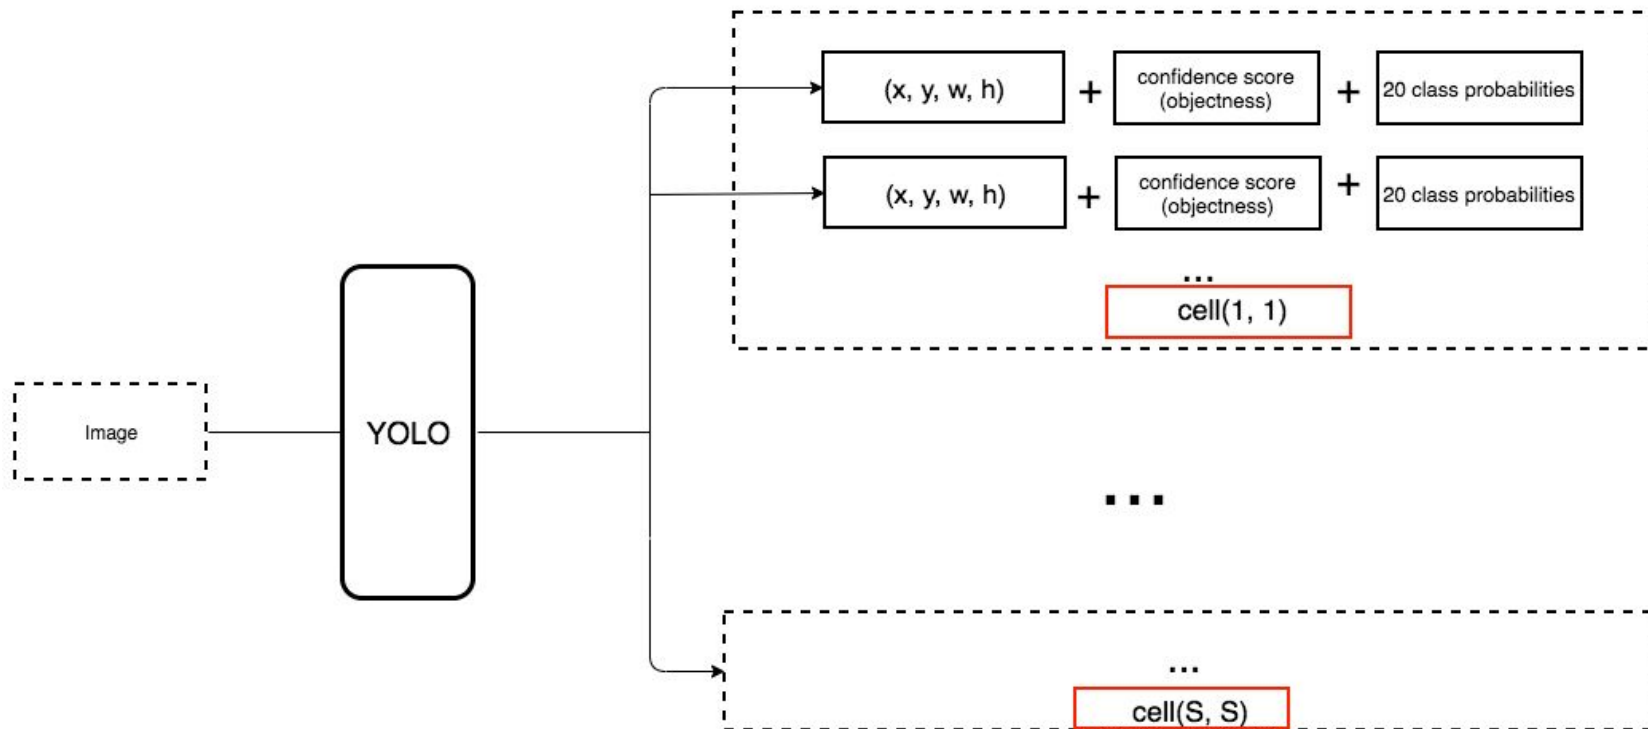
Class probability map

$Pr(\text{Class}_i | \text{object})$



Final detections


# YOLO






# YOLO

- Set a threshold beforehand in terms of IoU
- If proposed BBoxes has  $\text{IoU} > \text{thresh}$  leave this BB and remove other ones



Before non-max suppression




Non-Max  
Suppression



After non-max suppression




# SSD




# SSD


- Set of default BB of fixed size and ration
- Object of different sized is detected at different levels (depth)



(a) Image with GT boxes



(b)  $8 \times 8$  feature map



(c)  $4 \times 4$  feature map