

LLMs and Tools

Deep Learning

Aziz Temirkhanov
Lambda, HSE

Motivation

- **Knowledge** of LLM **constrained** to pretraining data

**GPT-5** Default  

The best model for coding and agentic tasks across domains

Compare Try in Playground

REASONING	SPEED	PRICE	INPUT	OUTPUT
		\$1.25 · \$10	   	   
Higher	Medium	Input · Output	Text, image	Text

GPT-5 is our flagship model for coding, reasoning, and agentic tasks across domains. Learn more in our [GPT-5 usage guide](#).

✦ 400,000 context window

↪ 128,000 max output tokens

 Sep 30, 2024 knowledge cutoff

 Reasoning token support

Motivation

- **Limited** context size

 **GPT-5** Default 

The best model for coding and agentic tasks across domains

REASONING



Higher

SPEED



Medium

PRICE

\$1.25 · \$10

Input · Output

INPUT



Text, image

OUTPUT



Text

Compare

Try in Playground

GPT-5 is our flagship model for coding, reasoning, and agentic tasks across domains. Learn more in our [GPT-5 usage guide](#).

✦ 400,000 context window


➞ 128,000 max output tokens

📅 Sep 30, 2024 knowledge cutoff

💡 Reasoning token support

Motivation


- LLM gets **distracted** by **useless** information



RAG

RAG = **R**etrieval-**A**ugmented **G**eneration

Idea. Augment prompt with **relevant** pieces of information.




Retrieve, Generate, Augment

1. **Retrieve** relevant document via similarity operation across the knowledge base




2. **Augment** prompt with retrieved information




3. **Generate** response




Create KB



Collect



Divide




Embed

Hyperparameters. Embedding size, chunk size, overlap between chunks

Candidate Retrieval

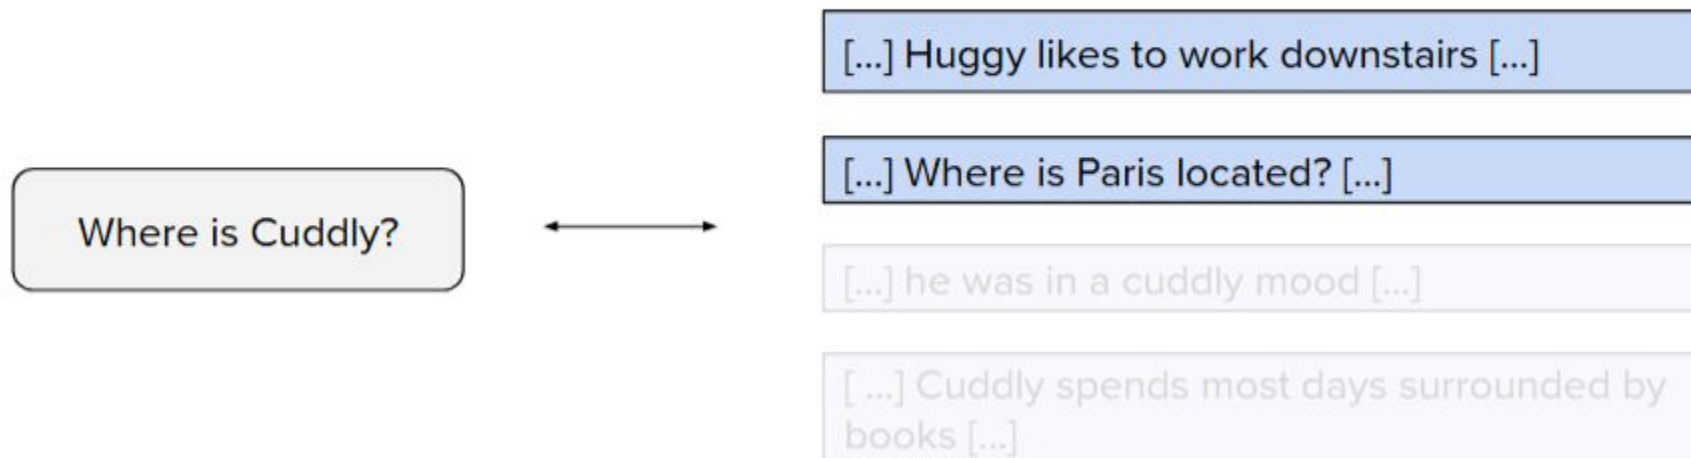
Objective. Select potentially-relevant candidates via search across knowledge base

Method 1. Semantic search using **embeddings-based similarity**



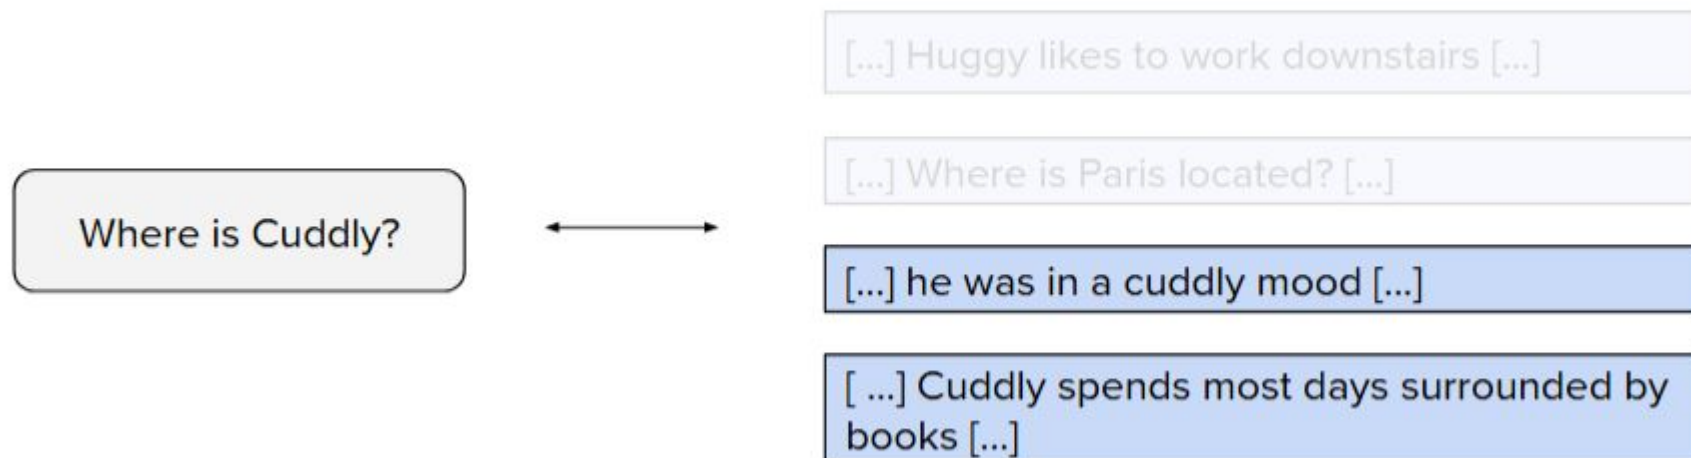
Candidate Retrieval

Method 1. Semantic search using **embeddings-based similarity**



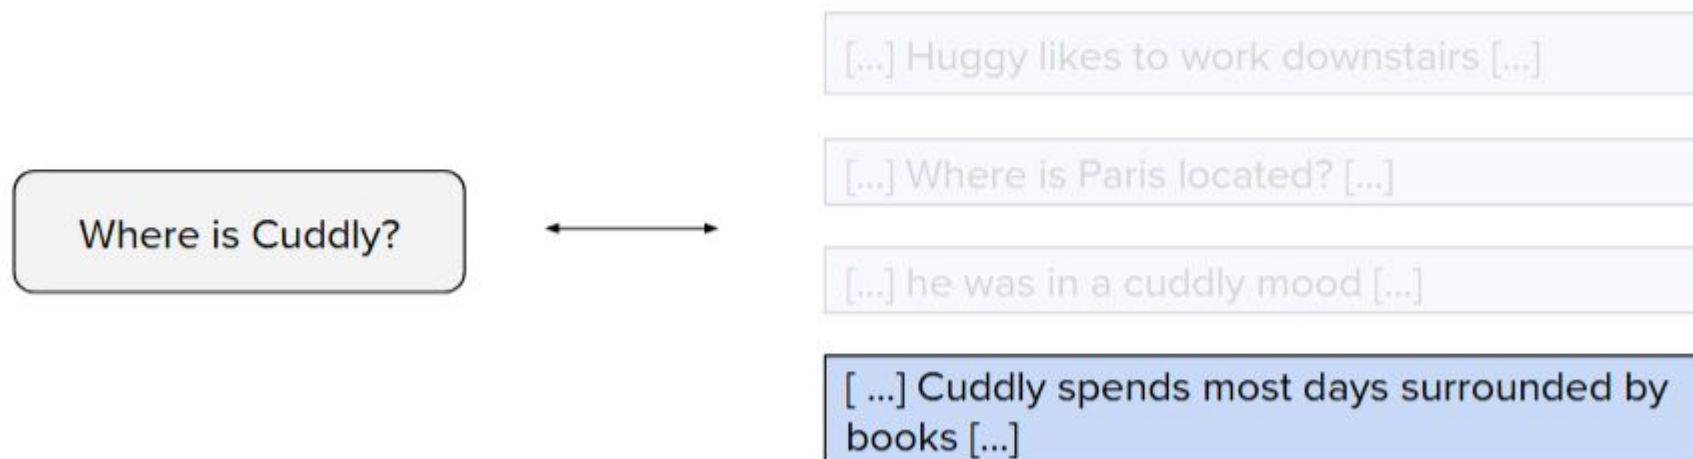
Candidate Retrieval

Method 2. Search based on keywords matching via **BM25**




Candidate Retrieval

Method 3. Search based on **hybrid combination** of semantics and BM25




Some More Methods

- Mitigate **discrepancy** in nature of embeddings



Some More Methods

- **Contextualize** document chunks




```
<document>
{WHOLE_DOCUMENT}
</document>
Here is the chunk we want to situate within the whole document:
{CHUNK_CONTENT}
```

Please give a short succinct context to situate this chunk within the overall document for the purposes of improving search retrieval of the chunk. Answer only with the succinct context and nothing else.


Ranking

Objective. Provide final relevance score using more sophisticated (re-)ranker

"cross-encoder"




Ranking



Ranking

Setup. Evaluate if **retrieved chunks** are **relevant**




Ranking

Setup. Evaluate if **retrieved chunks** are **relevant**

Ranking

- **N**ormalized **D**iscounted **C**umulative **G**ain at **k** (**NDCG@k**)



$$\text{DCG}@k = \sum_{i=1}^k \frac{\text{rel}_i}{\log_2(i+1)}$$

with $\text{rel}_i \in \{0, 1\}$

$$\text{NDCG}@k = \frac{\text{DCG}@k}{\text{IDCG}@k}$$


$\nearrow$
 $\text{DCG}@k$ if ranking was perfect

Ranking

Setup. Evaluate if **retrieved chunks** are **relevant**

Ranking

- **N**ormalized **D**iscounted **C**umulative **G**ain at **k** (**NDCG@k**)
- **R**eciprocal **R**ank at **k** (**RR@k**)



$$RR = \frac{1}{\text{rank}}$$


rank of the first relevant chunk

Ranking

Setup. Evaluate if **retrieved chunks** are **relevant**

Ranking

- **N**ormalized **D**iscounted **C**umulative **G**ain at **k** (**NDCG@k**)
- **R**eciprocal **R**ank at **k** (**RR@k**)
- **R**ecall at **k**




$$\text{Recall@}k = \frac{|\text{relevant in top } k|}{|\text{relevant}|}$$

Ranking

Setup. Evaluate if **retrieved chunks** are **relevant**

Ranking

- **N**ormalized **D**iscounted **C**umulative **G**ain at **k** (**NDCG@k**)
- **R**eciprocal **R**ank at **k** (**RR@k**)
- **R**ecall at **k**
- **P**recision at **k**



$$\text{Precision@}k = \frac{|\text{relevant in top } k|}{k}$$


Tools

Tools


- Web Search
- Math commuting
- Image Generation
- API Calls

Example

LLM as we've known it so far



LLM with tools



Example

1. Let **LLM** find **argument** for **relevant function** call



2. Make **function call**




3. Let **LLM** deduce conclusion **based on results**




How To Use Tools?

- Training
- Prompting



<function API> + <detailed explanation on how to use it>
Find a bear near me!



Summary

PROS

- LLMs just became way more useful!
- They can also interact with the real world
- Overcomes "knowledge cutoff" limitation


CONS

- More tools = decrease performance
- Finite context length: not scalable
- Many tools to define. Lots of work

MCP

MCP = Model Context Protocol

Idea. Connect tools/data to LLMs in a standard way



Agents

Idea


Traditional



Reasoning




Agent




ReAct

ReAct = **Reason** + **Act**



ReAct

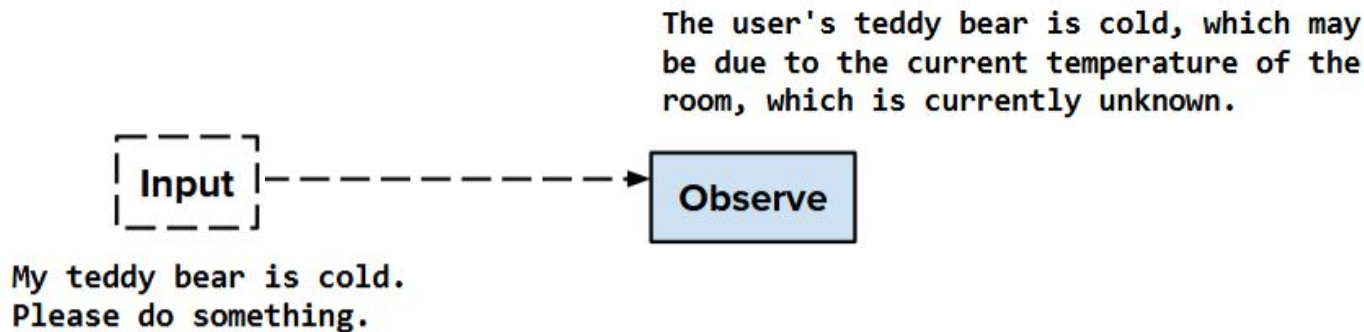


My teddy bear is cold.
Please do something.

Examples:

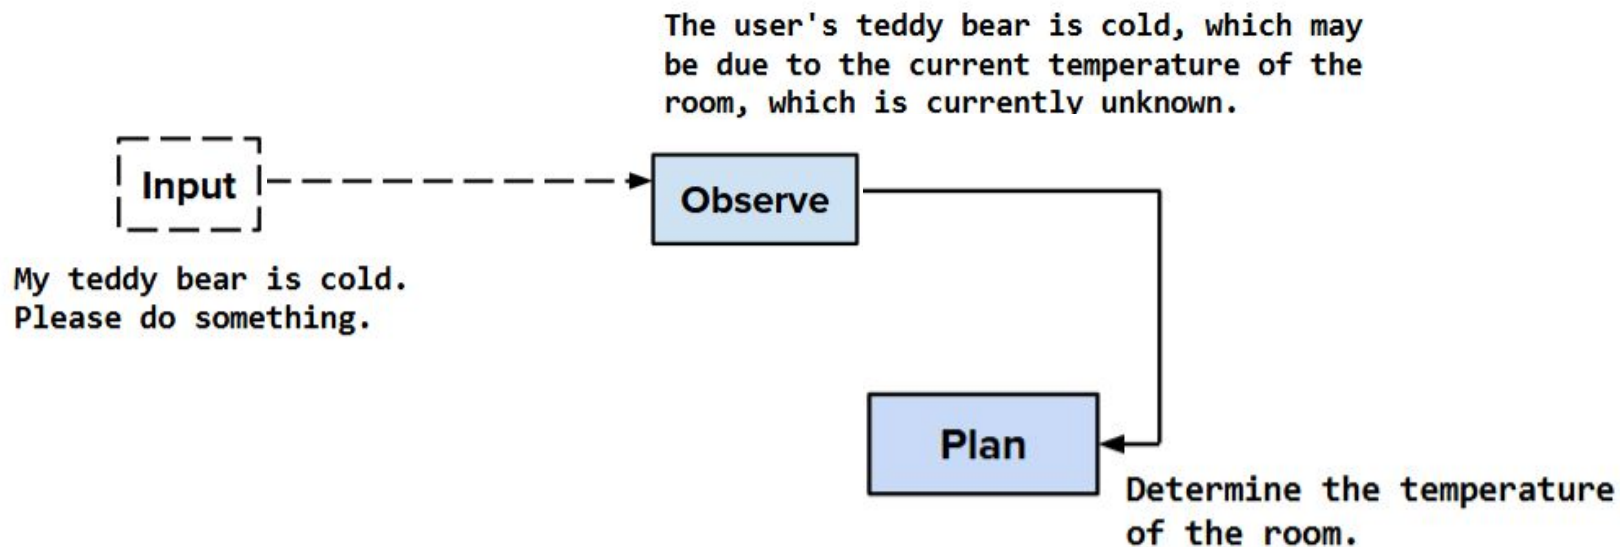
- Manually entered (e.g. user question)
- External event (e.g. metric going beyond a threshold)

ReAct




- **Synthesize** previous actions + explicitly **state** what is **currently known** including own knowledge
- **Reasoning**-heavy step to figure out **what is needed**

ReAct




Detail **what tasks** need to be accomplished and **what tools** to call

ReAct




- Perform an action via an **API**
- **Look** for info in a **database** of documents


ReAct



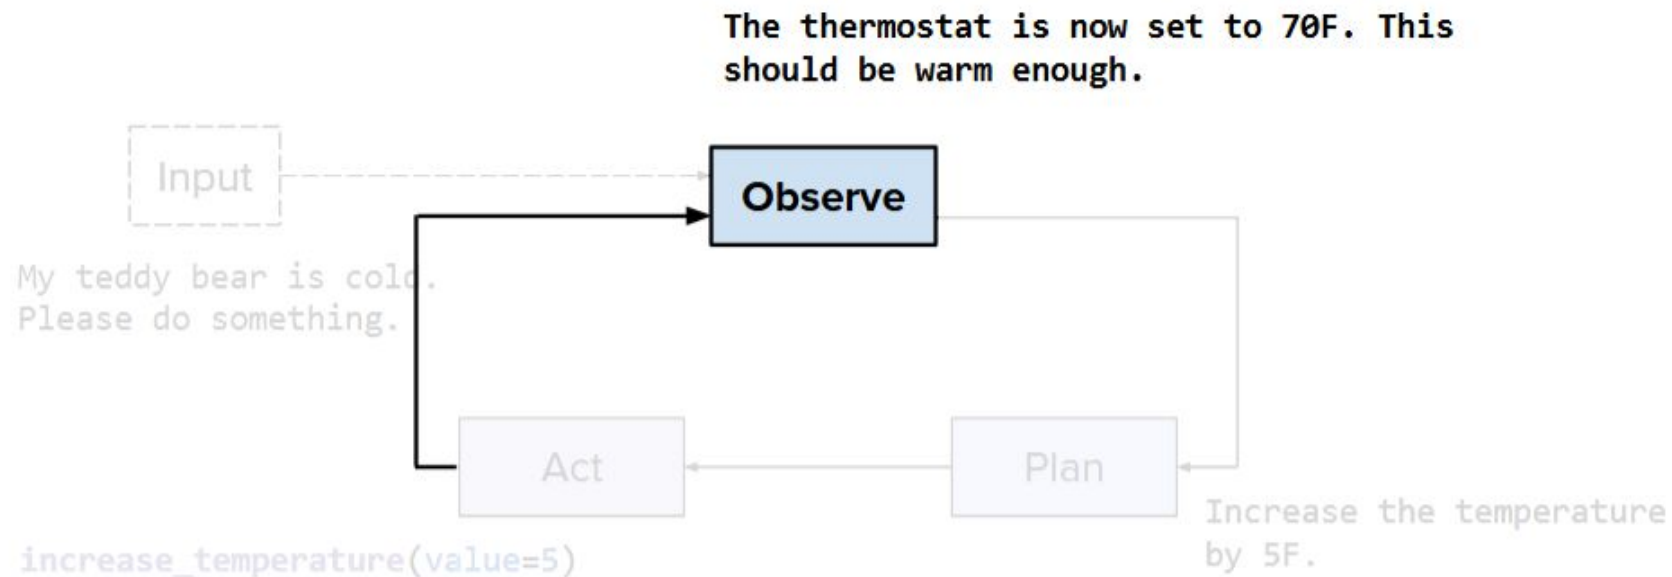
ReAct




ReAct




ReAct



ReAct




Agentic View



Agentic View

A2A = Agent2Agent



Safety

Risks.

- Potential for harm in the real world
- Example: data exfiltration

Remediations.

- Training steps
- Inference safeguards
- Benchmarks, e.g. **Agent-SafetyBench**

...very important topic!

Summary

- **Hallucination** is a (big) problem
- **Reasoning** abilities are a bottleneck
 - Finetuning helps, but hard
 - New capabilities are very welcome
- **Evaluation** is challenging
- Good to start simple, then iterate and progressively scale up
- Good to start with capable models, optimize on size later
- Transparency / observability helps with user trust and debuggability